

Rahul Rajaram

✉ rahul_rajjaram@outlook.com |

+1 (602) 552-2914

🔗 [Rahul2040](#) |

🌐 [rajaram-rahul](#) |

Irving, TX

PROFESSIONAL SUMMARY

Data Engineer with 2.5+ years of experience building cloud-native data pipelines on GCP, AWS, and Azure. Master's in Computer Science (Big Data Systems), currently supporting real-time analytics and pricing workflows at CVS Health. Skilled in PySpark, SQL, Airflow, and BigQuery, with a focus on pipeline optimization, legacy system migration, and scalable data solutions.

EDUCATION

Master of Science Computer Science (Big Data Systems)

Arizona State University | Tempe, AZ

Relevant Courses: Database Management System Implementation, Data Mining, Data Visualization, Distributed Data Systems, Cloud Computing, Multimedia and Web Data systems, Artificial Intelligence

Aug 2023 – May 2025

GPA: 3.80/4

Bachelor of Engineering in Electronics and Communication Engineering

Anna University | Chennai, TN, India

Relevant Courses: Data Structures and Algorithms, Cryptography and Network Security, Machine Learning Techniques

July 2017- May 2021

GPA: 8.21/10

SKILLS

Languages	Python, SQL, Bash
Tools & Technologies	Data Modeling, ETL, Data Warehousing, PySpark, Airflow, BigQuery, Snowflake, AWS S3, AWS DynamoDB, Jenkins, Docker, Git, Linux
Soft Skills	Problem-Solving, Communication, Team Collaboration

PROFESSIONAL EXPERIENCE

CVS health | Irving, TX, USA

Data Engineer

Jun. 2025 - Present

- Designed BigQuery data models and ETL pipelines for pricing, improving traceability, audit compliance, and analytical consistency across systems.
- Automated pricing workflows on GCP using PySpark, Airflow, and Jenkins, ensuring scalable, reliable deployments with clear SLA enforcement.
- Collaborated with data scientists and engineers to align models, metrics, and logging to product analytics and strategic pricing goals.
- Created reusable Python query generation engine supporting linear programming models, enhancing performance transparency and pricing logic adaptability.
- Implemented validation checks and dashboards to track data quality, SLA adherence, and operational reliability across all pricing pipelines.
- Led initiatives to enforce access controls and data lineage tracking, improving compliance and auditability of production systems and datasets.

Analytics Engineering Intern

May 2024 – Aug 2024 & Jan. 2025 – May 2025

- Migrated complex data pipelines from Teradata to **BigQuery** and **PySpark**, improving **scalability**, **performance**, and **cost-efficiency** across systems.
- Converted legacy SQL logic to **PySpark** with testing, ensuring exact match between old and new pipeline results under various conditions.
- Built and deployed **Airflow DAGs** with **Jenkins integration**, automating secure, reliable cloud workflows using GCP-native tools.
- Enabled **cluster parallelization** on DataProc, reducing execution time and costs by optimizing resource utilization and task scheduling.
- Partnered with product teams to expose cleaned datasets for **analytics**, **enabling timely insights** for pricing and operations.
- Enhanced metadata, **logging**, and **exception handling** to simplify monitoring and reduce time to diagnose issues in production workflows.

LTIMindtree | Chennai, TN, India

Senior Engineer IIOT

Aug. 2021 – June 2023

- Developed **ETL pipelines** for 1M+ **IoT** messages daily, optimizing transformation and schema for faster, cleaner business insights.
- Replaced costly **Databricks** workflows with **Snowpipe** and SQL, achieving **90% cost savings** without sacrificing performance or reliability.
- Created **React** and **Angular**-based interfaces with **PostgreSQL**, delivering interactive dashboards and tools for operations and analytics teams.
- Designed AWS-hosted web app for L&T Construction to track electricity usage and automate meter billing workflows.
- Coordinated across backend, frontend, and cloud teams, ensuring timely project delivery and effective cross-functional communication.
- Led internal AWS cloud training and mentored team on best practices for scalable architecture and secure data pipelines.

PROJECTS

Generic Drug Pricing

May 2024 - Present

- Built Python query generator engine to support LP pricing logic, enabling **modular**, **traceable**, and **auditable** pricing decision-making.
- Constructed **BigQuery** data pipelines orchestrated by **Airflow** and **Composer** to manage pricing updates with **high availability** and **compliance**.
- Designed scalable workflows on **DataProc** clusters to handle **complex model logic** and **deliver timely** pricing outputs.
- Reduced manual interventions by **automating** execution and validations through **Jenkins**, improving deployment **reliability** and regulatory readiness.
- Collaborated with product, legal, and analytics teams to ensure pricing logic aligned with business goals and compliance standards.
- Drove implementation of monitoring dashboards and failure alerts, enhancing transparency, debuggability, and operational visibility for stakeholders.

Columnar Database Design

Jan. 2024 – May 2024

- Developed a custom column-store database on Java-based Minibase to optimize analytical queries involving **grouping**, **filtering**, and **projections**.
- Extended query engine to support columnar operations with **indexing**, boosting performance in **analytical** workloads versus row-store storage.
- Designed and tested storage format to improve access speed for large-scale aggregations and minimize read latency.
- Implemented support for selection, projection, and group-by, benchmarking against traditional formats to demonstrate analytical advantages.
- Strengthened debugging and Java system-level skills while troubleshooting performance bottlenecks in query execution.
- Gained deep insights into database internals and architectural tradeoffs involved in analytical database engine design.

Connected Batteries

Nov. 2022- Jun. 2023

- Integrated IoT sensors on trucks to capture telemetry and GPS data for **real-time fleet and battery health monitoring**.
- Built data pipelines using **Azure IoT Hub**, **Eventhub**, and **Snowflake** to process streaming telemetry at scale.
- Created stored procedures for data cleaning and transformation, improving consistency and downstream usability of raw data.
- Coordinated with product and cloud teams to deliver reliable, end-to-end pipeline for battery monitoring analytics.
- Delivered business dashboards with fleet-level analytics, aiding predictive maintenance and reducing vehicle downtime.
- Ensured production readiness with exception handling, retry logic, and system health metrics in Snowflake-based workflows.